

# SQL (Part 4)

## Database Modification Statements and Transactions in SQL

**Instructor: Shel Finkelstein**

- *Database Modification Reference: A First Course in Database Systems, 3<sup>rd</sup> edition, Chapter 6.5*
- *Transactions Reference: Transactions Reference: A First Course in Database Systems, 3<sup>rd</sup> edition, Chapter 6.6 – 6.7*

# Important Notices

- Lab3 assignment is due on **Sunday, May 19**, by 11:59pm.
  - 3 weeks to do Lab3 because of Midterm.
  - Lab2 was/will be discussed at Lab Sections.
  - Your solution should be submitted via Canvas as a zip file.
    - Canvas is used for both Lab submission and grading.
    - Late Lab Assignments will not be accepted.
    - Be sure that you post the correct file!
  - Load file for Lab3 has been posted to Piazza.
    - You must use load file to do Lab3.
    - Helps with testing, but we won't post query solutions.
- The **third** Gradiance Assignment, "CMPS 182 Spring 2019 #3", was/will be posted on Monday, April 29, and is due **Monday, May 6**, by 11:59pm.

# Important Notices

- **Reminder:** Midterm is on **Wednesday, May 8.**
  - **Includes material up to and including this Lecture**, but not Lecture 7.
  - No make-ups, no early Midterms, no late Midterms ... and no devices.
  - You may bring a **single two-sided 8.5" x 11" sheet of paper** with as much info written (or printed) on it as you can fit and read unassisted.
    - No sharing of these sheets will be permitted.
  - “Practice Midterm” from Spring 2017 has been posted on Piazza under Resources → Exams.
    - Solution will be posted there next week ... but take it yourself first.
    - Some questions are on topics we haven't covered yet.
  - Hope that all requests for DRC accommodation have been submitted.
  - Piazza announcement will describe required seating pattern for Midterm.
- See [Small Group Tutoring website](#) for LSS Tutoring with [Chandler Hawkins](#).

# Database Modification Statements

**Instructor: Shel Finkelstein**

*Database Modification Reference:  
A First Course in Database Systems,  
3<sup>rd</sup> edition, Chapter 6.5*

# Database Modification Statements

- SQL statements for:
  - *Inserting some* tuples into a relation
  - *Deleting some* tuples from a relation
  - *Updating* values of some columns of some existing tuples
- INSERT, DELETE, and UPDATE are referred to as *modification* operations.
  - They are Data Manipulation Language (DML) statements, as is SELECT.
- Modification operations change the *state* of the database.
  - They do not return a collection of rows or other values.
  - They may return errors/error codes.

# Insert Statement with Values

```
INSERT INTO R(A1, ..., An)  
VALUES (v1, ..., vn);
```

- A tuple (v<sub>1</sub>, ..., v<sub>n</sub>) is inserted into the relation R, where attribute A<sub>i</sub> = v<sub>i</sub> and default values (perhaps NULL) are entered for all missing attributes.

```
INSERT INTO StarsIn(movieTitle, movieYear, starName)  
VALUES ('The Maltese Falcon', 1942, 'Sydney Greenstreet');
```

- The tuple ('The Maltese Falcon', 1942, 'Sydney Greenstreet') will be added to the relation StarsIn.

```
INSERT INTO StarsIn  
VALUES ('The Maltese Falcon', 1942, 'Sydney Greenstreet');
```

# INSERT Statement with Subquery

Movies(title, year, length, genre, studioName, producerC#)  
Studio(name, address, presC#)

```
INSERT INTO Studio(name)
  SELECT DISTINCT studioName
  FROM Movies
  WHERE studioName NOT IN
    (SELECT name
     FROM Studio);
```

Why DISTINCT?

- Add to the relation Studio all the names that appear in the studioName column of Movies but do not already occur in the names in the Studio relation.

# INSERT Statement with Subquery and Constants

Movies(title, year, length, genre, studioName, producerC#)  
Disney2018Movies(title, length)

```
INSERT INTO Movies(title, year, length, studioName)
  SELECT dm.title, 2018, dm.length, 'Disney'
  FROM Disney2018Movies dm;
```

Puts the Disney2018Movies into the Movies table, setting year and studioName values using constants.

What will be the genre and producerC# values for the inserted movies?



# DELETE Statement

```
DELETE FROM R  
WHERE <condition>;
```

```
DELETE FROM StarsIn  
WHERE movieTitle = 'The Maltese Falcon' AND  
      movieYear = 1942 AND  
      starName = 'Sydney Greenstreet';
```

- The tuple ('The Maltese Falcon', 1942, 'Sydney Greenstreet') will be deleted from the relation StarsIn.
- What if we wanted to delete tuples from StarsIn for all movies starring Sydney Greenstreet?

# More DELETE Examples

```
DELETE FROM MovieExec  
WHERE netWorth < 10000000;
```

- Deletes all movie executives whose net worth is less than 10 million dollars.

```
DELETE FROM MovieExec  
WHERE cert# IN  
  (SELECT m.producerC#  
   FROM Movies m, StarsIn s  
   WHERE m.title = s.movieTitle AND m.year = s.movieYear  
        AND s.starName = 'Sydney Greenstreet');
```

- Deletes all movie executives who produced movies starring Sydney Greenstreet

# DELETE: Careful

What does:

```
DELETE FROM MovieExec;
```

without a WHERE clause do?

Answer: Deletes all the tuples from MovieExec!!!

# UPDATE Statement

```
UPDATE R
  SET <new-value-assignments>
  WHERE <condition>;
```

- <new-value-assignment> :-  
    <attribute> = <expression>, ..., <attribute> = <expression>

```
UPDATE Employees
  SET salary = 85000, dept = 'SALES'
  WHERE SSnum='123456789';
```

```
UPDATE Employees
  SET salary = 25000
  WHERE salary IS NULL;
```

```
UPDATE Employees
  SET salary = salary * 1.1
  WHERE salary > 100000;
```

# UPDATE with Subquery

UPDATE R

SET <new-value-assignments>

WHERE <condition>;

- <new-value-assignment>:-  
    <attribute> = <expression>, ..., <attribute> = <expression>

UPDATE MovieExec

SET name = 'Pres.' || name

WHERE cert# IN ( SELECT presC# FROM Studio );

- 2<sup>nd</sup> line: concatenates the string 'Pres.' with name.

# Semantics of Modifications

- Database modification statements are completely evaluated on the old state of the database, producing a new state of the database.
  - What does this statement do? Is it deterministic or not?

```
UPDATE MovieExec e
  SET e.NetWorth = 6M
  WHERE NOT EXISTS (SELECT * FROM MovieExec e2
                    WHERE e2.Networth = 6M);
```

Should write 6000000,. This is clearer (but it's not legal SQL).

MovieExec

name	address	cert#	netWorth
S. Spielberg	X	38120	3M
G. Lucas	Y	43918	4M
W. Disney	Z	65271	5M


# Transactions in SQL

**Instructor: Shel Finkelstein**

*Transactions Reference:  
A First Course in Database Systems,  
3<sup>rd</sup> edition, Chapter 6.6 – 6.7*

# One-Statement-At-a-Time Semantics

- So far, we have learnt how to query and modify the database.
- SQL statements posed to the database system were executed one at a time, retrieving data or changing the data in the database.





# Transactions

- Applications such as web services, banking, airline reservations demand high throughput on operations performed on the database.
  - Manage hundreds of sales transactions every second.
  - Transactions often involve multiple SQL statements.
  - Database are transformed to new state based on (multiple statement) transactions, not just single SQL statements.
- It's possible for two operations to simultaneously affect the same bank account or flight, e.g. two spouses doing banking transactions, or an automatic deposit during a withdrawal, or two people reserving the same seat.
  - These “concurrent” operations must be handled carefully.

# ACID Transactions

- Atomicity
- Consistency
- Isolation
- Durability

# Simple Example of What Could Go Wrong

Flights(fltNo, fltDate, seatNo, seatStatus, purchaser)

- Customer1 issues the following query via a web application.

```
SELECT seatNo
```

```
FROM Flights
```

```
WHERE fltNo=123 AND fltDate=DATE '2012-12-25'  
      AND seatStatus='available';
```

- Customer1 inspects the results and selects a seat, say 22A.

```
UPDATE Flights
```

```
SET seatStatus='occupied', purchaser='Customer1'
```

```
WHERE fltNo=123 AND fltDate= DATE '2012-12-25' AND seatNo='22A';
```

# Simple Example of What Could Go Wrong (continued)

- Customer2 is also looking at the same flight on the same day simultaneously and decides to choose seat 22A as well.
- Operations of query and update statements:

<< Draw on Board >>

- Both customers believe that they have reserved seat 22A.
- Problem: Each SQL statement of both users is executed correctly, but the overall result is not correct.
- However, a DBMS can provide the illusion that the actions of Customer1 and Customer2 are executed *serially* (i.e., one at a time, with no overlap).
  - **Serializability**

# Another Example of What Could Go Wrong, Even with a Single User

Accounts(acctNo, balance)

- User1 wants to transfer \$100 from an account with acctNo=123 to an account with acctNo=456.

- 1. Subtract \$100 from the account with acctNo=123**

UPDATE Accounts

SET balance = balance – 100

WHERE acctNo=123;

- 2. Add \$100 to the account with acctNo=456**

UPDATE Accounts

SET balance = balance + 100

WHERE acctNo=456;

- What if application or database fails after step 1, but before step 2?

# Atomicity

- Failure (e.g., network failure, power failure etc.) could occur after step 1.
  - If this happens, money has been withdrawn from account 123 ...
  - ... but not not deposited into account 456.
- The DBMS should provide mechanisms to ensure that groups of operations are executed **atomically**.
  - That is, either **all** the operations in the group are executed to completion or **none** of the operations are executed.
  - All-or-nothing, no in-between

# Transactions

- A *transaction* is a group of operations that should be executed atomically, all-or-nothing.
- Operations of a transaction can be interleaved with operations of other transactions.
- However, with an “isolation level” called *serializability*, the **illusion** is given that every transaction is executed one-by-one, in a serial order.
  - The DBMS will execute each transaction in its entirety or not at all, “without transactions interfering with each other”.

# Transactions (cont'd)

- START TRANSACTION or BEGIN TRANSACTION (can be implicit)
  - Marks the beginning of a transaction, followed by one or more SQL statements.
- COMMIT
  - Ends the transaction. All changes to the database caused by the SQL statements within the transaction are committed (i.e., they are permanently there--**Durability**) and visible in the database.
  - All changes become visible at once (atomically).
  - Before commit, changes to the database caused by the SQL statements are visible to this transaction, but are not visible to other transactions.
- ROLLBACK
  - Causes the transaction to abort or terminate. Any changes made by SQL statements within the transaction are undone (“rolled back”).



# Example Using Informal Syntax

BEGIN TRANSACTION

<SQL statement to check whether bank account 123 has  $\geq$  \$100>

If there is no account 123, then ROLLBACK;

If account 123 has  $<$  \$100, then ROLLBACK;

<SQL statement to withdraw \$100 from account 123>

<SQL statement to add \$100 to account 456>

If there is no account 456, then ROLLBACK;

COMMIT;

- Scenario 1: Suppose bank account 123 has \$50.
- Scenario 2: Bank account 123 has \$200, bank account 456 has \$400.
- Scenario 3: Bank account 123 has \$200, bank account 456 has \$400, failure after withdrawing \$100 from account 123.
- Scenario 4: Bank account 123 has \$200, bank account 456 has \$400, failure after depositing \$100 to account 456, but before COMMIT

# Read-Only Transactions

- In the previous examples, each transaction involved a read, then a write.
- If a transaction has only read operations, it is less likely to impact serializability.
- SET TRANSACTION READ ONLY;
  - Stated *before* the transaction begins.
  - Tells the SQL system that the next transaction is read-only.
  - SQL may take advantage of this knowledge to parallelize many read-only transactions.
- SET TRANSACTION READ WRITE;
  - Tells SQL that the next transaction may write data, in addition to read.
  - Default option if not specified; often (usually) not specified.

# Dirty Reads (Read Uncommitted)

- *Dirty data* refers to data that is written by a transaction but has not yet been committed by the transaction.
- *A dirty read* refers to the read of dirty data written by another transaction.
- Consider the following transaction T that transfers an amount of money (\$X) from one account to another:
  1. Add \$X to Account 2.
  2. Test if Account 1 has \$X.
    - a) If there is insufficient money, remove \$X from Account 2.
    - b) Otherwise, subtract \$X from Account 1.

# Dirty Reads (cont'd)

- Transaction T1: Transfers \$150 from A1 to A2.
- Transaction T2: Transfers \$250 from A2 to A3.
- Initially: A1: \$100, A2: \$200, A3: \$300.
- What might be the (unexpected, unwanted) result with Dirty Reads if execution of T1 and T2 happens to interlace in a certain way?

<< To discuss, and write on board >>

# Should Transactions Allow Dirty Reads?

- **Allow** Dirty Reads
  - More parallelism between transactions.
  - But may cause serious problems, as previous example shows.
- **Don't Allow** Dirty Reads
  - Less parallelism, more time is spent on waiting for other transactions to commit or rollback.
  - More overhead in the DBMS to prevent dirty reads.
  - Cleaner semantics.

# Isolation levels

```
SET TRANSACTION READ WRITE  
ISOLATION LEVEL READ UNCOMMITTED;
```

- First line: The transaction may write data (that's the default).
- Second line: The transaction can run with isolation level "Read Uncommitted", allowing Dirty Reads.
- Default Isolation Level depends on system.
  - Most systems run with READ COMMITTED or SNAPSHOT ISOLATION.

# Other Isolation Levels

- SET TRANSACTION ISOLATION LEVEL READ COMMITTED;
  - Only clean (committed) reads, no dirty reads.
  - But you might read data committed by *different* transactions.
    - You might not even get the same value even when you read same data a second time during a single transaction!
- SET TRANSACTION ISOLATION LEVEL REPEATABLE READ;
  - Repeated queries of a tuple during a transaction will retrieve the same value, even if its value was changed by another transaction.
    - But *different* data reads might return values that were committed by *different* transactions at different times.
  - Also, a second scan of a range (e.g., salary>10000) may return “phantoms” not originally present in the scan..
    - Phantoms are tuples newly inserted while the transaction is running.
- SET TRANSACTION ISOLATION LEVEL SERIALIZABLE;

# Isolation Levels

Isolation Level	dirty reads	non-repeatable reads	phantoms
READ UNCOMMITTED	Y	Y	Y
READ COMMITTED	N	Y	Y
REPEATABLE READ	N	N	Y
SERIALIZABLE	N	N	N



# Snapshot Isolation (SI)

- SI and Read Committed are most commonly used Isolation Levels.
  - Better performance (response time, throughput) than Serializability
- Transaction reads data as it existed when transaction began (repeatable).
  - As usual, transaction also sees its own updates
- Conflicts on Writes are avoided; equivalent of Serializable on Writes.
  - ... but not on Read/Write interactions between transactions
- Example: Two transactions that are running under **Serializability** change both A and B. If both Commit, then one ran logically after the other.
- **SI Example:** A is supposed to be less than B. Initially A is 0 and B is 100.
  - T1 reads original A and B values, and changes A to 60.
  - T2 reads original A and B values, and changes B to 20.
  - Transactions T1 and T2 both maintained the consistency condition “ $A < B$ ” ... but what are the final values of A and B?
  - Could this happen with Serializability?