

# Constraints and Triggers

**Instructor: Shel Finkelstein**

*Reference:*

*A First Course in Database Systems,  
3<sup>rd</sup> edition, Chapter 7 (but not section 7.4)*

# Important Notices

- Lab3 assignment is due on **Sunday, May 19**, by 11:59pm.
  - 3 weeks to do Lab3 because of Midterm.
  - Lab2 was/will be discussed at Lab Sections.
  - Some material in this Lecture (Lecture 8) is used in Lab3 (Foreign Key and General Constraints).
  - Your solution should be submitted via Canvas as a zip file.
    - Canvas is used for both Lab submission and grading.
    - Late Lab Assignments will not be accepted.
    - Be sure that you post the correct file!
  - Load file for Lab3 has been posted to Piazza.
    - You must use load file to do Lab3, and original data must be used at multiple stages of Lab3.
    - Load data helps with testing, but we won't post query solutions.

# Important Notices

- **Reminder:** Midterm was on Wednesday, May 8.
  - Exam and Answers have been posted to Piazza under Resources→Exams
  - Answers were/will be discussed in class on Friday, May 10.
  - Midterms will be/were returned during last 10 minutes of class on Monday, May 13.
  - If you have questions about the grading of the Midterm, please speak to your TA or to me soon!
- See [Small Group Tutoring website](#) for LSS Tutoring with [Chandler Hawkins](#).

# Constraints and Triggers

- A *constraint* is a relationship among data elements that the DBMS is required to enforce.
  - *Example*: key constraints.
- *Triggers/Rules* are only executed when a specified condition occurs, e.g., insertion of a tuple.
  - Easier to implement than complex constraints.

# Kinds of Constraints

- Keys/Unique constraints
- Foreign-key, or referential-integrity constraints
- Value-based constraints
  - Constrain values of a particular attribute
- Tuple-based constraints
  - Relationship among components of tuple
- Assertions
  - Any SQL boolean expression (not implemented in most relational DBMS, not discussed in this lecture)

# Review: Single-Attribute Keys

- Place PRIMARY KEY or UNIQUE after the type in the declaration of the attribute.
- Example:

```
CREATE TABLE Beers (  
    name    CHAR(20)  UNIQUE,  
    manf    CHAR(20)  
);
```

# Review: Multi-Attribute Key

- The bar and beer together are the key for Sells:

```
CREATE TABLE Sells (  
    bar      CHAR(20),  
    beer     VARCHAR(20),  
    price    REAL,  
    PRIMARY KEY (bar, beer)  
);
```

# Review: NULL

```
CREATE TABLE Sells (  
    bar          CHAR(20),  
    beer         VARCHAR(20),  
    price REAL NOT NULL,  
    PRIMARY KEY (bar, beer)  
);
```

If the CREATE statement didn't include NOT NULL for price:

```
ALTER TABLE Sells ALTER COLUMN price SET NOT NULL;
```

```
ALTER TABLE Sells ALTER COLUMN price DROP NOT NULL;
```



# Foreign Keys

- Values appearing in attributes of one relation must also appear together in specific attributes of another relation.
- **Example:**
  - In **Sells(bar, beer, price)**, we might expect that a **beer** value also appears in Beers.name (the name column of the Beers table, the primary key for that table).
- Like a link/pointer, but based on **value**.

# Example: Sales Record Schema

Products(productID, productName, manuf, normalPrice, discount)

Customers(customerID, custName, address, joinDate, amountOwed,  
lastPaidDate, status)

Stores(storeID, storeName, region, address, manager)

Days(dayDate, category)

Sales(productID, customerID, storeID, dayDate, paidPrice, quantity)

Payments(customerID, custName, paidDate, amountPaid, cleared)

Foreign Key Constraints that might make sense include:

- Every customerID in Sales is a customerID Primary Key in Customers
- Every productID in Sales is a productID Primary Key in Products
- Every storeID in Sales is a storeID Primary Key in Stores
- Every customerID in payments is a customerID Primary Key in Customers

# Expressing Foreign Keys

- Use keyword REFERENCES, either:
  1. After an attribute (for one-attribute keys)
  2. As an element of the schema:  
FOREIGN KEY (<list of attributes>  
REFERENCES <relation> (<attributes>)
- Referenced attributes must be declared as either PRIMARY KEY or UNIQUE.
- (Why?)

## Example: With Attribute

```
CREATE TABLE Beers (  
    name CHAR(20) PRIMARY KEY,  
    manf CHAR(20) );  
  
CREATE TABLE Sells (  
    bar CHAR(20),  
    beer CHAR(20) REFERENCES Beers(name),  
    price REAL );
```

## Example: As Schema Element

```
CREATE TABLE Beers (  
    name CHAR(20) PRIMARY KEY,  
    manf CHAR(20) );  
  
CREATE TABLE Sells (  
    bar CHAR(20),  
    beer CHAR(20),  
    price REAL,  
    FOREIGN KEY(beer) REFERENCES  
        Beers(name) );
```

## Example: Adding Foreign Key

```
CREATE TABLE Beers (  
    name CHAR(20) PRIMARY KEY,  
    manf CHAR(20) );  
  
CREATE TABLE Sells (  
    bar CHAR(20),  
    beer CHAR(20),  
    price REAL );  
  
ALTER TABLE Sells  
    ADD FOREIGN KEY (beer)  
    REFERENCES Beers(name);
```

# Enforcing Foreign-Key Constraints (Referential Integrity, RI)

- If there is a foreign-key constraint from *referring relation*  $R$  to *referenced relation*  $S$ , then violations may occur two ways:
  1. An insert or update to  $R$  introduces values that are not found in  $S$ , or
  2. A deletion or update to  $S$  causes some tuples of  $R$  to “dangle”, referencing a value that no longer exists

# Actions Taken --- (1)

- **Example:** suppose  $R = \text{Sells}$ ,  $S = \text{Beers}$ .
  - That is, Sells refers to Beers
- An insert or update to **Sells** that introduces a nonexistent beer must be rejected.
- A deletion or update to **Beers** that removes a beer value found in some tuples of **Sells** can be handled in one of three ways (next slide).



# Actions Taken --- (2)

1. *Default* : Reject the modification.
2. *Cascade* : Make the same changes in Sells.
  - **Deleted beer**: delete corresponding Sells tuples.
  - **Updated beer**: also change value in Sells ...
    - ... so that Sells.beer has the same new value as Beers.name
3. *Set NULL* : Change Sells.beer to NULL.

# Example: Cascade

- Upon Delete of the Bud tuple from Beers:
  - Delete all tuples from Sells that have beer = 'Bud'
- Upon Update of the Bud tuple by changing 'Bud' to 'Budweiser':
  - Change all Sells tuples that have beer = 'Bud' to have beer = 'Budweiser'

## Example: Set NULL

- Upon Delete of the Bud tuple from Beers:
  - Change all tuples of Sells that have beer = 'Bud' to have beer = NULL.
- Upon Update of the Bud tuple, changing 'Bud' to 'Budweiser':
  - Make the same change to tuples of Sells that have beer='Bud' as for deletion (making beer=NULL).

# Choosing a Referential Integrity Policy

- When we declare a foreign key, we may choose policies SET NULL or CASCADE independently for deletions and updates.
- Follow the foreign-key declaration with:  
`ON [UPDATE, DELETE][SET NULL, CASCADE]`
- Two such clauses may be used, one for UPDATE and one for DELETE
- Otherwise, the DEFAULT (Reject) is used.

## Example: Setting Policy

```
CREATE TABLE Sells (  
    bar    CHAR(20),  
    beer   CHAR(20),  
    price  REAL,  
    FOREIGN KEY (beer)  
        REFERENCES Beers (name)  
        ON DELETE SET NULL  
        ON UPDATE CASCADE  
);
```

# Attribute-Based Check

- Constraint on the value of a particular attribute.
- CHECK(<condition>) may be added to the declaration for the attribute.
  - Condition must evaluate to TRUE or UNKNOWN; can't be FALSE.
- The condition may refer to the attribute of the relation that is being checked.
- But for the condition to reference any other tuples or relations, a subquery must be used.
  - Note: PostgreSQL does not support CHECK with subquery. Sigh.

## Example: Attribute-Based Check

```
CREATE TABLE Sells (  
    bar CHAR(20),  
    beer CHAR(20) CHECK ( beer IN  
        (SELECT name FROM Beers) ),  
    price REAL CHECK ( price <= 5.00 )  
);
```

# Example: Named Constraints

```
CREATE TABLE Sells (  
    bar CHAR(20),  
    beer CHAR(20) CHECK ( beer IN  
        (SELECT name FROM Beers)),  
    price REAL  
        CONSTRAINT price_is_cheap  
        CHECK ( price <= 5.00 )  
);
```

```
ALTER TABLE Sells DROP CONSTRAINT price_is_cheap;  
ALTER TABLE Sells ADD CONSTRAINT price_is_cheap  
    CHECK ( price <= 5.00 );
```



# Timing of Attribute-Based Checks

- Attribute-based checks are performed only when a value for that attribute is inserted or updated.
  - **Example:** `CHECK (price <= 5.00)` checks every new price and rejects the modification (for that tuple) if the price in Sells is more than \$5.
  - **Example:** `CHECK (beer IN (SELECT name FROM Beers))` is not checked if a beer is deleted from Beers (unlike foreign-keys).

# Tuple-Based Checks

- CHECK (<condition>) may be added as a relation-schema element.
- The condition may refer to any attribute of the relation (same tuple).
- But for the condition to reference any other tuples or relations, a subquery must be used.
  - Condition is checked **only** on INSERT or UPDATE into relation that has the CHECK.

## Example: Tuple-Based Check

- Only Joe's Bar can sell beer for more than \$5:

```
CREATE TABLE Sells (  
  bar      CHAR(20),  
  beer     CHAR(20),  
  price    REAL,  
  CHECK (bar = 'Joe's Bar' OR  
         price <= 5.00)  
);
```

# Assertions

- These are database-schema elements, like relations or views.
- Defined by:  
    CREATE ASSERTION <name>  
        CHECK (<condition>);
- Condition may refer to any relation or attribute in the database schema.
- (**Not implemented** in most Relational DBMS because they're too complicated and expensive!)

# Triggers: Motivation

- Assertions are powerful, but the DBMS often can't tell when they need to be checked ...
  - ... and they're probably not implemented by the DBMS.
- Attribute- and tuple-based checks are checked at known times, but they are not that powerful.
- Triggers let the user (often the DBA) decide when to check for any condition.

# Event-Condition-Action Rules

- Another name for “trigger” is an *ECA Rule*, or *Event-Condition-Action* Rule
- *Event* : typically a type of database modification, e.g., “insert on Sells”
- *Condition* : Any SQL boolean-valued expression
- *Action* : Any SQL statements

# Preliminary Example: A Trigger

- Instead of using a foreign-key constraint and rejecting insertions into `Sells(bar, beer, price)` when beer doesn't appear as a name in the relation `Beers(name, manf)`, a trigger can add that beer to `Beers`, with a NULL manufacturer.

# Example: Trigger Definition

```
CREATE TRIGGER BeerTrig
```

```
AFTER INSERT ON Sells
```

The event

```
REFERENCING NEW ROW AS NewTuple  
FOR EACH ROW
```

```
WHEN (NewTuple.beer NOT IN  
      (SELECT name FROM Beers))
```

The condition

```
INSERT INTO Beers(name)  
VALUES(NewTuple.beer);
```

The action



# CREATE TRIGGER

- Either:

`CREATE TRIGGER <name>`

- Or:

`CREATE OR REPLACE TRIGGER <name>`

- Useful if there is a trigger with that name and you want to modify the trigger.

# Options: The Event

- AFTER INSERT can be BEFORE INSERT.
  - Also, can be INSTEAD OF, if the relation is a view.
    - A clever way to execute view modifications is to have triggers translate them to appropriate modifications on the base tables.
- INSERT can be DELETE or UPDATE.
  - And UPDATE can be UPDATE OF a particular attribute.

# Options: FOR EACH ROW

- Triggers are either “row-level” or “statement-level.”
- FOR EACH ROW indicates row-level; its absence indicates statement-level.
- *Row level triggers* : Execute once for each modified tuple.
- *Statement-level triggers* : Execute once for a SQL statement, regardless of how many tuples are modified.

# Options: REFERENCING

- INSERT statements imply a new tuple (for row-level) or new table (for statement-level).
  - The “table” is the set of inserted tuples.
- DELETE implies an old tuple or table.
- UPDATE implies both.
- Refer to these by  
[NEW OLD] [TUPLE TABLE] AS <name>

# Options: The Condition

- Any boolean-valued condition.
- Evaluated on the database as it would exist before or after the triggering event, depending on whether BEFORE or AFTER is used.
  - But always before any triggered changes take effect.
- Access the new/old tuple/table through the names in the REFERENCING clause.

# Options: The Action

- There can be more than one SQL statement in the action.
  - Surround by BEGIN . . . END if there is more than one.
- But queries make no sense in an action, so we are really limited to modifications.

# Another Example

- Using `Sells(bar, beer, price)` and a unary relation `RipoffBars(bar)`, maintain a list of RipoffBars that raise the price of some beer by more than \$1.

# The Trigger

CREATE TRIGGER PriceTrig

AFTER UPDATE OF price ON Sells

REFERENCING

OLD ROW AS ooo

NEW ROW AS nnn

FOR EACH ROW

WHEN(nnn.price > ooo.price + 1.00)

INSERT INTO RipoffBars  
VALUES(nnn.bar);

The event –  
only changes  
to prices

Updates let us  
talk about old  
and new tuples

Condition:  
a raise in  
price > \$1

We need to consider  
each price change

When the price change  
is great enough, add  
the bar to RipoffBars